


RISC-V Instruction Set Architecture and Simulators

October 15, 2019

Institute of Embedded Systems (E-13)
Hamburg University of Technology

Instruction Set Architecture



- The **Interface** between Software and Hardware
- The **Programmers view** of the processor
- **Hides the details** of processor implementation
- Enables a **wide variety of microarchitecture** implementations without compiler customization

What is RISC-V?

- Fifth generation of RISC design from UC Berkeley
- A high-quality, [license-free, royalty-free RISC ISA](#) specification ([here](#))
- > 25 RISC-V CPU Designs
- Supported by [growing shared software ecosystem](#)
- Appropriate for all levels of computing system, [from microcontrollers to supercomputers](#)
- Standard maintained by [non-profit RISC-V Foundation](#) (Foundation Members (60+))

Note: The RISC-V ISA manual include useful information about the made design decisions

RISC-V instruction length encoding

32-Bit RISC-V: RV32I (Base Integer Instructions)

64-Bit RISC-V: RV64I (Base Integer Instructions)

128-Bit RISC-V: RV128I (Base Integer Instructions)

xxxxxxxxxxxxxxxxaa

 16-bit ($aa \neq 11$)

xxxxxxxxxxxxxxxxxx	xxxxxxxxxxxbbb11
--------------------	------------------

 32-bit ($bbb \neq 111$)

≪≪≪ XXXX	xxxxxxxxxxxxxxxxxx	xxxxxxxxxxx011111
----------	--------------------	-------------------

 48-bit

≪≪≪ XXXX	xxxxxxxxxxxxxxxxxx	xxxxxxxxxxx011111
----------	--------------------	-------------------

 64-bit

≪≪≪ XXXX	xxxxxxxxxxxxxxxxxx	xxxxxnnnn111111
----------	--------------------	-----------------

 $(80 + 16 * nnnn)$ -bit, $nnnn \neq 1111$

≪≪≪ XXXX	xxxxxxxxxxxxxxxxxx	xxxxx1111111111
----------	--------------------	-----------------

 Reserved for ≥ 320 -bits

RISC-V ISA Extensions

M: Integer Multiply and Division *

F: Single-Precision Floating-Point (32-bit) *

D: Double-Precision Floating-Point (64-bit) *

Q: Quad-Precision Floating-Point (128-bit) *

L: Decimal-Precision Floating-Point

C: Compressed Instructions (16-bit encoding)

expands to 32-bit instruction (2-address format + 32 Regs | 3-address format + 8 Regs),
smallest Instruction Format on SPECint2006

B: Bit Manipulation Instructions

J: Dynamically Translated Languages

T: Transactional Memory

P: Packed SIMD Instructions *

V: Vector Operations *

N: User-Level Interrupts *

G: General-Purpose (I+M+A+F+D)

The TinyRV ISA is:

- A **subset** of the 32-bit RISC-V ISA suitable for teaching
- More specifically, the TinyRV ISA is a subset of the **RV32IM ISA**
- suitable for **running simple C programs**
- suitable for a **FPGA** implementation
- From an assembler programmer's perspective, the **RV32IM ISA is very similar to the MIPS32 ISA**: Many assembler programs are executable after register renaming.

- Tiny RISC-V supports only
 - Words / int: 32-bit
 - Half-words / short int: 16-bit
 - Bytes / char: 8-bit
 - No floating-point
- General Purpose Registers:
 - 31 general-purpose registers x1-x31, each register is 32 bits wide
 - Register x0 is hardwired to the constant zero
 - Registers have no type (Operation determines how register contents are interpreted)
 - TinyRV uses the same calling convention and symbolic register names as RISC-V

TinyRV Registers and its usage by the compiler

Mnemonic	Reg Nr.	Description	Saver
zero	x0	constant value 0	
ra	x1	return address	Caller
sp	x2	stack pointer	Callee
gp	x3	global pointer	
tp	x4	thread pointer	
t0 - t2	x5 - x7	temporary registers	Caller
s0 / fp	x8	saved register or frame pointer	Callee
s1	x9	saved register	Callee
a0 - a1	x10 - x11	function arguments / return values	Caller
a2 - a7	x12 - x17	function arguments	Caller
s2 - s11	x18 - x27	saved registers	Callee
t3 - t6	x28 - x31	temporary registers	Caller

TinyRV Memory

- TinyRV supports **only a 1MB virtual memory address space** from 0x00000000 to 0x000fffff
- **Memory addresses are in bytes**, not words, so **word addresses are 4 bytes apart** ($PC = PC + 4$)
- The memory system is of **little endianness** type:
Endianness specifies the order of bytes in the target register when we load a word from memory:

- Assume:

Data	A	B	C	D
Address	0x0	0x1	0x2	0x3
- If we load a four-byte word from address 0x0, there are two options for the destination register:
 - 0xABCD (big endian)
 - **0xDCBA** (little endian).

There is **no significant benefit** of one system over the other


TinyRV Instruction Encoding

- The TinyRV ISA uses the [same instruction encoding as RISC-V](#)
- Each instruction has a specific instruction type

31	30	25	24	21	20	19	15	14	12	11	8	7	6	0			
funct7				rs2			rs1		funct3		rd			opcode		R-type	
imm[11]		imm[10:5]		imm[4:1]		imm[0]		rs1		funct3		rd			opcode		I-type
imm[11]		imm[10:5]		rs2			rs1		funct3		imm[4:1]		imm[0]		opcode		S-type
imm[12]		imm[10:5]		rs2			rs1		funct3		imm[4:1]		imm[11]		opcode		B-type
imm[31]		imm[30:20]					imm[19:15]		imm[14:12]		rd			opcode		U-type	
imm[20]		imm[10:5]		imm[4:1]		imm[11]		imm[19:15]		imm[14:12]		rd			opcode		J-type

TinyRV Immediate Formats

- Note that in RISC-V all immediates are always sign extended
- Sign-bit for the immediate is always in bit 31 of the instruction.
- The following diagrams illustrate how to create a 32-bit immediate from each of the five immediate formats.



TinyRV Integer Register-Register Operations

Inst	Name	FMT	Opcode	F3	F7	Description (C)	Note
add	ADD	R	0x3	0x0	0x00	$rd = rs1 + rs2$	
sub	SUB	R		0x0	0x20	$rd = rs1 - rs2$	
xor	XOR	R		0x4	0x00	$rd = rs1 \wedge rs2$	
or	OR	R		0x6	0x00	$rd = rs1 \mid rs2$	
and	AND	R		0x7	0x00	$rd = rs1 \& rs2$	
sll	Shift Left Logical	R		0x1	0x00	$rd = rs1 \ll rs2$	
srl	Shift Right Logical	R		0x2	0x00	$rd = rs1 \gg rs2$	
sra	Shift Right Arith	R		0x3	0x20	$rd = rs1 \gg rs2$	msb-ext.
slt	Set Less Than	R	0x33	0x2		$rd = (rs1 < rs2)?1:0$	
sltu	Set Less Than (U)	R		0x3		$rd = (rs1 < rs2)?1:0$	zero-ext.
mul	multiply	R		0x0	0x01	$rd = (rs1 * rs2)[31:0]$	

TinyRV Integer Register-Immediate Operations

Inst	Name	FMT	Opcode	F3	F7	Description (C)	Note
addi	ADD Immediate	I	0x13	0x0	0x00	$rd = rs1 + imm$	
xori	XOR Immediate	I		0x4	0x00	$rd = rs1 \wedge imm$	
ori	OR Immediate	I		0x6	0x00	$rd = rs1 imm$	
andi	AND Immediate	I		0x7	0x00	$rd = rs1 \& imm$	
slli	Shift Left Logical Imm	I		0x1	0x00	$rd = rs1 \ll imm$	
srlr	Shift Right Logical Imm	I		0x1	0x00	$rd = rs1 \gg imm$	
srai	Shift Right Arith Imm	I		0x3	0x20	$rd = rs1 \ggg imm$	msb-ext.
slti	Set Less Than Imm	I		0x2		$rd = (rs1 < imm)?1:0$	
sltiu	Set Less Than Imm (U)	I		0x3		$rd = (rs1 < imm)?1:0$	zero-ext.
lui	Load Upper Imm	U	0x37			$rd = imm \ll 12$	
auipc	Add Upper Imm to PC	U	0x17			$rd = PC + (imm \ll 12)$	

TinyRV Load and Store Operations

Inst	Name	FMT	Opcode	F3	F7	Description (C)	Note
lb	Load Byte	I	0x3	0x0		$rd = M[rs1+imm][0:7]$	
lh	Load Half	I		0x1		$rd = M[rs1+imm][0:15]$	
lw	Load Word	I		0x2		$rd = M[rs1+imm][0:31]$	
lbu	Load Byte (U)	I		0x4		$rd = M[rs1+imm][0:7]$	zero-ext.
lhu	Load Half (U)	I		0x5		$rd = M[rs1+imm][0:15]$	zero-ext.
sb	Store Byte	S	0x23	0x0		$M[rs1+imm][0:7] = rs2[0:7]$	
sh	Store Half	S		0x1		$M[rs1+imm][0:15] = rs2[0:15]$	
sw	Store Word	S		0x2		$M[rs1+imm][0:31] = rs2[0:31]$	
csrr	Atomic Read CSR	I	0x73	0x2		$rd = CSR[csr]$	csr=imm
csrw	Atomic Write CSR			0x1		$CSR[csr] = rs1$	csr=imm

Control Transfer Operations

Inst	Name	FMT	Opcode	F3	F7	Description (C)	Note
beq	Branch ==	B	0x63	0x0		if (rs1 == rs2) PC += imm	
bne	Branch !=	B		0x1		if (rs1 != rs2) PC += imm	
blt	Branch <	B		0x4		if (rs1 < rs2) PC += imm	
bge	Branch <=	B		0x5		if (rs1 >= rs2) PC += imm	
bltu	Branch < (Unsigned)	B		0x6		if (rs1 < rs2) PC += imm	
bgeu	Branch >= (Unsigned)	B		0x7		if (rs1 >= rs2) PC += imm	
jal	Jump And Link	J	0xCF			rd = PC+4; PC += imm	zero-ext.
jalr	Jump And Link Reg	I	0x37	0x0		rd = PC+4; PC = rs1	zero-ext.
ecall	Environment Call	I	0x73	0x0	0x00	Transfer control to OS a0 = 1: print value of a1	imm: 0x000
ebreak	Environment Break	I	0x73	0x0	0x00	Transfer control to debug	imm: 0x001

TinyRV Privileged ISA?

- No distinction between user and privileged mode
 - (Using the terminology in the RISC-V vol 2 ISA manual, TinyRV only supports M-mode)
- Reset Vector
 - RISC-V specifies two potential reset vectors: one at a low address, and one at a high address.
 - TinyRV uses the low address reset vector at 0x00000200.
 - This is where assembly tests should reside as well as user code in TinyRV.

Control/Status Registers (CSRs)

TinyRV includes five non-standard CSRs:

mng2proc: 0xFC0

Used to [communicate data from the manager to the processor](#).

This register has register-mapped FIFO-dequeue semantics meaning reading the register essentially dequeues the data from the head of a FIFO.

Reading the register will stall if the FIFO has no valid data.

Writing the register is undefined.

proc2mng2: 0x7c0

Used to [communicate data from the processor to the manager](#).

This register has register-mapped FIFO-enqueue semantics meaning writing the register essentially enqueues the data on the tail of a FIFO.

Writing the register will stall if the FIFO is not ready.

Reading the register is undefined.

Control/Status Registers (CSRs)

stats_en: 0x7c1

Used to **enable or disable the statistics tracking** feature of the processor (i.e., counting cycles and instructions)

numcores: 0xfc1

Used to store the **number of cores present in a multi-core system**

Writing the register is undefined.

coreid: 0xf14

Used to communicate the **core id in a multi-core system**

Writing the register is undefined

Address Translation

- TinyRV only supports the most basic form of address translation:
 - `logical address = physical address`
- In the RISC-V vol 2 ISA manual this is called a `Mbare` addressing environment.

Pseudo Instructions

Pseudoinstruction	Base Instruction(s)	Meaning
la rd, symbol	auipc rd, symbol[31:12] addi rd, rd, symbol[11:0]	Load address
l{b h w d} rd, symbol	auipc rd, symbol[31:12] l{b h w d} rd, symbol[11:0](rd)	Load global
s{b h w d} rd, symbol, rt	auipc rt, symbol[31:12] s{b h w d} rd, symbol[11:0](rt)	Store global
nop	addi x0, x0, 0	No operation
li rd, immediate	Myriad sequences	Load immediate
mv rd, rs	addi rd, rs, 0	Copy register
not rd, rs	xori rd, rs, -1	One's complement
neg rd, rs	sub rd, x0, rs	Two's complement
negw rd, rs	subw rd, x0, rs	Two's complement word
sext.w rd, rs	addiw rd, rs, 0	Sign extend word
seqz rd, rs	sltiu rd, rs, 1	Set if = zero
snez rd, rs	sltu rd, x0, rs	Set if != zero
sltz rd, rs	slt rd, rs, x0	Set if < zero
sgtz rd, rs	slt rd, x0, rs	Set if > zero

Code examples: Transfers between Memory and Register

- Word data transfers

```
lw x10,12(x13)  — load  word  
sw x10,40(x13)  — store word
```

- x13 – base register;
- 12 – offset in bytes; Offset must be a constant known at assembly time

Code examples: Long immediates

How to set 0xDEADBEEF?

```
lui  x10, 0xDEADB    # x10 = 0xDEADB000  
addi x10, x10, 0xEEF  # x10 = 0xDEADAEFF
```

Immediates are always sign-extended, if top bit is set, will subtract -1 from upper 20 bits

Solution:

```
lui  x10, 0xDEADC    # x10 = 0xDEADC000  
addi x10, x10, 0xEEF  # x10 = 0xDEADBEEF
```

Assembler pseudo-op handles all of this:

```
li  x10, 0xDEADBEEF  # Creates two instructions
```

Code examples: Translation of `if` statement

f: x10, g: x11, h: x12, i: x13, j: x14

```
if ( i == j )                                bne x13,x14,Exit
    f = g + h;                               add x10,x11,x12
                                           Exit:
```

```
if ( i == j )                                bne x13,x14,Exit
    f = g + h;                               add x10,x11,x12
else                                          j Exit
    f = g - h;                               Else: sub x10,x11,x12
                                           Exit:
```

Code examples: Translation of `switch` statement

```
switch (amount) {  
    case 20:      case20:  addi x1, x0, 20      # x2:amount, x3:fee  
        fee = 2;    bne x2, x1, case50          # x1 = 20  
                    # amount == 20? if not  
                    # skip to case50  
                    addi x3, x0, 2             # if so, fee = 2  
                    j done                     # and break out of case  
    break;  
  
    case 50:      case50:  addi x1, x0, 50      # x1 = 50  
        fee = 3;    bne x2, x1, case100         # amount == 50? if not,  
                    # skip to case100  
                    addi x3, x0, 3             # if so, fee = 3  
                    j done                     # and break out of case  
    break;  
  
    case 100:     case100: addi x1, x0, 100     # x1 = 100  
        fee = 5;    bne x2, x1, default        # amount == 100? if not,  
                    # skip to default  
                    addi x3, x0, 5             # if so, fee = 5  
                    j done                     # and break out of case  
    break;  
  
    default:     default: add x3,x0,x0         # fee = 0  
        fee = 0;    done:  
}  
}
```


Code examples: Translation of a `while` loop

```
int pow = 1;
int x = 0;

while (pow != 128) {
    pow = pow * 2;
    x = x + 1;
}

addi x1, x0, 1      # x1 = pow,  pow = 1
addi x2, x0, 0      # x = 0,  x2 = x,
addi x3, x0, 128    # x3 = 128 for comparison
while: beq x1, x3, done # if pow == 128, goto done
        slli x1, x1, 1  # pow = pow * 2 by shiftleft
        addi x2, x2, 1  # x = x + 1
        j while
done: mv a1, x2      # exit while loop
      li a0, 1
      ecall          # print x2
```

Code examples: Translation of a for loop

```
int A[5] = {20,30,10,40,50};  
  
int sum = 0;  
for (int i=0; i<5; i++)  
    sum += A[i];
```

```
len: .data                                # data segment  
      .word 5                            # Length of A[]  
A:    .word 20, 30, 10, 40, 50  
      .text  
      la x9, A                          # x9 = &A[0]  
      add x10, x0, x0                    # sum = 0  
      add x11, x0, x0                    # i = 0  
  
Loop: lw x12, 0(x9)                      # x12 = A[i]  
      add x10,x10,x12                    # sum +=  
      addi x9, x9,4                      # &A[i++]  
      addi x11,x11,1                     # i++  
      la x13, len                        # x13 = 5  
      lw x13 0(x13)  
      blt x11,x13,Loop
```

Code examples: Terminal Output

```
.data
msg:  .asciiz "\nhello world"
pi:   .float 3.1415927

.text
start: li a0, 1          # ecall code: print integer
      li a1, 0xa        # integer to print
      ecall
      li a0, 4          # ecall code: print string
      la a1, msg        # null-terminated string address
      ecall
      li a0, 11         # ecall code: print character
      li a1, 'a'        # character to print
      ecall
      li a0, 10         # ecall code: stops running
      ecall
```

- RISC-V doesn't have a concept of a screen, but you can use the `ecall` instruction, a system call, to output characters to the console (which might be a serial port or a pseudo tty).
- <https://riscvsim.com/ecalls/>

Code examples: Translation of array accesses

```
.data                                # data segment

len:    .word 5                      # Length of arrays
val:    .word 20, 30, 10, 40, 50

.text

la      x5, val                     # x5 = &val[0]
lw      x4, 0(x5)                   # x4 = val[0]
slli    x4, x4, 3                   # x4 = x4 << 3 = x4 * 8
sw      x4, 0(x5)                   # val[0] = x4
lw      x4, 4(x5)                   # x4 = val[1]
slli    x4, x4, 3                   # x4 = x4 << 3 = x4 * 8
sw      x4, 4(x5)                   # val[1] = x4
```

Seven Fundamental Steps in Calling a Function

- ① Put parameters in a place where the function can access them
- ② Transfer control to function
- ③ Acquire (local) storage resources needed for function
- ④ Perform desired task of the function
- ⑤ Put result value in a place where calling code can access it
- ⑥ Restore any registers you used
- ⑦ Return control to point of origin, since a function can be called from several points in a program

RISC-V Function Call

- Registers faster than memory, so use them
- a0 – a7 (x10 - x17) : eight argument registers to pass parameters and two return values (a0 - a1)
- ra : one return address register to return to the point of origin (x1)

```
• 1000      mv a0,s0           # x = a
  1004      mv a1,s1           # y = b
  1008      addi ra,zero,1016    # ra =1016
  1012      j sum # jump to sum
  1016      ...                # next instruction ...
  2000  sum : add a 0,a0,a1
  2004      jr ra              # jump register instruction
```

- **Assembler shorthand: ret = jr ra**
- sum might be called by many places, so we can't return to a fixed place
- **1008 jal sum # ra =1012,goto sum**

(really should be laj “link and jump”)

Stack: Memory for register values before function call

Where are old register values saved to restore them after function call?

- Need a place to save old values before call function, restore them when return, and delete
- Ideal is a stack, a last-in-first-out (LIFO) queue with its two operations:
 - Push: placing data onto stack
 - Pop: removing data from stack
- Stack is located in data memory, so need register to point to it
 - **sp is the stack pointer in RISC-V**
- Convention is grow stack down from high to low addresses
 - Push decrements `sp`
 - Pop increments `sp`

Function example


```
int Leaf (int g, int h, int i, int j) {  
    int f;  
    f = (g+h)-(i+j);  
    return f;  
}
```

- Parameter variables g, h, i, and j in argument registers a0, a1, a2, and a3, and f in s0
- Assume need one temporary register s1


- Leaf:

```
addi sp, sp, -8      # adjust stack for 2 items  
sw    s1, 4 ( sp )   # save s1 for use afterwards  
sw    s0, 0 ( sp )   # save s0 for use afterwards  
add   s0, a0, a1      # s0 = g + h  
add   s1, a2, a3      # s1 = i + j  
sub   s0, s0, s1      # return value (g + h) - (i + j)  
lw    s0, 0 ( sp )   # restore register s0 for caller  
lw    s1, 4 ( sp )   # restore register s1 for caller  
addi  sp, sp, 8      # adjust stack to delete 2 items  
jr    ra             # jump back to calling routine
```



Stack before, during, after Function



Before call



During call



After call

Nested Procedures: Need to save the return address of the caller

```
int sumSquare(int x, int y) {  
    return mult (x,x) + y ;  
}
```

- Something called `sumSquare` , now `sumSquare` is calling `mult`
- So there's a value in `ra` that `sumSquare` wants to jump back to, but this will be overwritten by the call to `mult`
- Need to save `sumSquare` return address before call to `mult`
- In general, may need to save some other info in addition to `ra`.

Optimized Function Convention

To reduce expensive loads and stores from spilling and restoring registers, RISC-V function - calling convention divides registers into two categories:

① Preserved across function call

- Caller can rely on values being unchanged
- `sp`, `gp`, `tp`, “saved registers” `s0-s11` (`s0` is also `fp`)

② Not preserved across function call

- Caller cannot rely on values being unchanged
- Argument/return registers `a0-a7`, `ra`, “temporary registers” `t0-t6`, callee can use temporary registers without saving and restoring them

Allocating Space on Stack

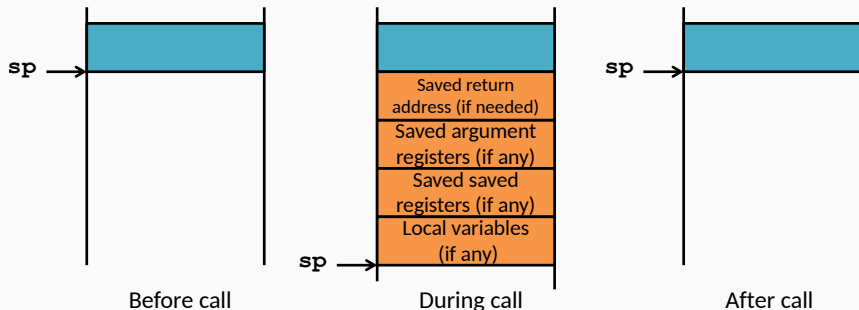
C has two storage classes:

Automatic variables are local to function and discarded when function exits

Static variables exist across exits from and entries to procedures

- Use stack for automatic (local) variables that don't fit in registers

Procedure frame / activation record : segment of stack with saved registers and local variables



Memory Layout

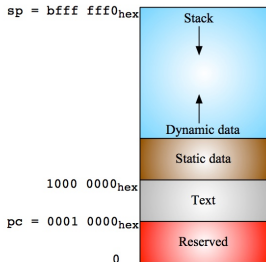
When a C program is run, there are three important memory areas allocated:

Static : Variables declared once per program, e.g. C globals

Heap : Variables declared dynamically via `malloc`

Stack : Space to be used by procedure during execution; this is where we can save register values

- Stack must be aligned on 16-byte boundary (not true in examples above)
- Global pointer (`gp`) points to static, RV32 `gp` = 1000_0000 hex
- **Heap above static for data structures that grow and shrink**



Instruction Set Simulator

Should be

- Accurate (cycle/timing or at least only functional)
- Fast to simulate entire systems ($>100\text{MIPS}$)
- expandable for new interfaces to the outer world
- customizable to new instructions
- Allows Multicore simulations



Instruction-Set Simulators (Venus)

The screenshot displays the Venus RISC-V simulator interface. At the top, there are tabs for 'Editor' and 'Simulator', with 'Simulator' being the active tab. Below the tabs are control buttons: 'Run' (highlighted in green), 'Step', 'Prev', 'Reset', and 'Dump'.

The main area is divided into three columns: 'Machine Code', 'Basic Code', and 'Original Code'. The 'Basic Code' column contains assembly instructions, and the 'Original Code' column contains their corresponding RISC-V assembly code. The first instruction is highlighted in green.

Machine Code	Basic Code	Original Code
0x000000ef	jai x1 0	main: jai ioutt
0x000000ef	jai x0 0	end: j end
0x00100113	addi x2 x0 1	for_init: addi x2, x0, 1
0x10000397	auipc x7 65536	lw x7, len
0xff43a393	lw x7 -12(x7)	lw x7, len
0x00100093	addi x1 x0 1	addi x1, x0, 1
0x401383b3	sub x7 x7 x1	sub x7, x7, x1
0x10000417	auipc x8 65536	la x8, val # pseudo instr.: load address: reg[8] = addr(val[0]);
0xfe840413	addi x8 x8 -24	la x8, val # pseudo instr.: load address: reg[8] = addr(val[0]);
0x00400493	addi x9 x0 4	addi x9, x0, 4
0x0423c463	blt x7 x2 72	for_loop: bgt x2, x7, end_for
0x401101b3	sub x3 x2 x1	sub x3, x2, x1

Below the code table is a 'console output' area.

On the right side, there is a 'Registers Memory' panel. It shows a list of registers (x0 to x15) and their current values. The 'x0' register is highlighted in green. Below the registers, there is a 'Display Settings' dropdown menu set to 'Hex'.

Venus (<http://www.kvakil.me/venus/>): Simple online simulator from Berkeley to learn RISC-V assembly language

Venus: Features

- RV32IM ISA
- Single-step debugging with undo feature
- Breakpoint debugging
- View machine code and original instructions side-by-side
- Several ecalls: including `print` and `sbrk`
- Memory visualization
- Integrated CodeMirror editor with syntax highlighting
- In order to assemble code, simply click on the "Simulator" tab

Venus: (Pseudo) Instructions

- Supported Instructions

At the moment, venus supports most of the RV32IM standard. `ebreak`, `fence`, `fence.i` and cycle counter instructions are not currently supported

- Pseudo Instructions

Venus currently supports all standard RISC-V pseudoinstructions as long as they translate into a series of supported instructions

Venus: Assembler directives

The following assembler directives are supported:

Directive	Effect
<code>.data</code>	Store subsequent items in the [static segment]
<code>.text</code>	Store subsequent instructions in the [text segment]
<code>.byte</code>	Store listed values as 8-bit bytes
<code>.ascii</code>	Store subsequent string in the data segment and add null-terminator
<code>.word</code>	Store listed values as unaligned 32-bit words
<code>.global</code>	Makes the given label global

Venus: Environmental Calls

- To use an environmental call, load the ID into register a0, and load any arguments into a1 - a7
- Any return values will be stored in argument registers

The following environmental calls are currently supported:

ID (a0)	Name	Description
1	print_int	prints integer in a1
4	print_string	prints the null-terminated string whose address is in a1
9	sbrk	allocates a1 bytes on the heap, returns pointer to start in a0
10	exit	ends the program
11	print_character	prints ASCII character in a1
17	exit2	ends the program with return code in a1

As an example, the following code prints the integer 42 to the console:

```
addi a0 x0 1  # print_int ecall
addi a1 x0 42  # integer 42
ecall
```

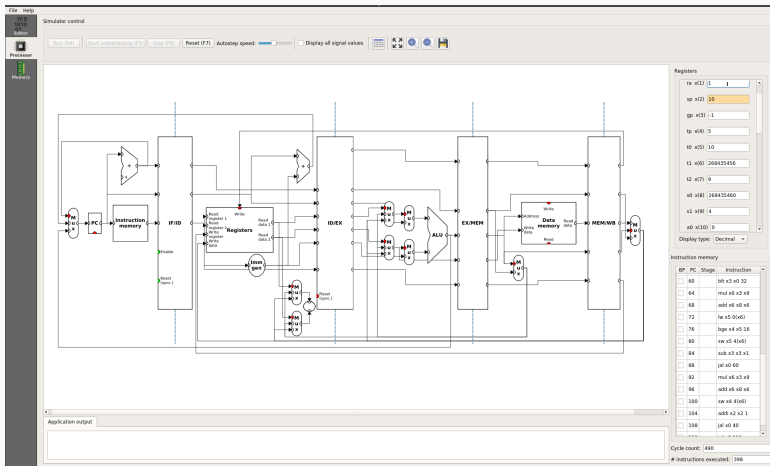
Memory Segments

Memory is divided into four different segments: text, static, heap and stack

Segment	Starting Address
text	0x0000_0000
static	0x1000_0000
heap	0x1000_8000
stack*	0x7fff_fff0

* Unlike other memory segments, the stack grows towards smaller addresses.

Instruction-Set Simulators (Ripes)



eebcRipes (<https://github.com/mortbopet/Ripes>): 5-stage pipelined RISC-V simulator